


Spring 2020 - Knowledge Discovery in Data at Scale Technologies - Homework #3

by Almaz Murzabekov



Homework description

Linear Regression

Write a program which illustrates simple linear regression (or a more general variant of linear regression) and implements accumulation of canonical information.

- a)** For some fixed parameters a and b (or, in a more general case, a_1, \dots, a_m) generate a sequence of observations (x_i, y_i) :

$$y_i = f(x_i) + \varepsilon_i,$$

where

$$f(x) = a + bx$$

or

$$f(x) = a_1 + a_2 x + a_3 x^2 + \dots + a_m x^{(m-1)}$$

ε_i are i.i.d. with zero mean and $E\varepsilon_i = \sigma$. Values x_i can be generated randomly with some mean and variance.

- b)** Accumulate canonical information, i.e., at each step, when a new observation (x_i, y_i) is produced, update canonical information

- c)** Illustrate the real function $f(x)$ and its estimate $\widehat{f(x)}$.

- d)** Illustrate $\widehat{\text{Var}(f(x))}$, assuming that σ^2 is known.

- e)** Illustrate $\widehat{\text{Var}(f(x))}$, assuming that σ^2 is NOT known.

In your report present the source code and a few (around 3) nice graphs showing estimations for “small”, “intermediate”, and “large” number of observations.

God damn! How to solve it? I hate the math! :(me too

Relax, don't panic!

Function used in the demo - polynomial:

$$y_i = 1 + 1 \cdot x_i - 1 \cdot x_i^2 + 0.2 \cdot x_i^3 + \varepsilon_i$$

Data:

$$(x_i, y_i), i = 1, \dots, n$$

Canonical information:

$$(T, v, V, n)$$

Elementary information:

$$(T_i, v_i, V_i, n_i)$$

Update:

$$(T, v, V, n) + (T_i, v_i, V_i, n_i) = (T + T_i, v + v_i, V + V_i, n + n_i)$$

Estimate $f(x)$:

$$(T, v, V, n) * x =>$$

- $\widehat{f(x)} = F_x T^{-1} v$
- $Var(\widehat{f(x)}) = \sigma^2 F_x T^{-1} F_x^T$
- $Var(\widehat{f(x)}) = \frac{V - v^T T^{-1} v}{n - m} * F_x T^{-1} F_x^T$

Now, as you see it becomes simple! Just define the functions and plot three pictures?!

Ok, but I don't got it! What is the T, v, V , and n ?

Hmm... Let's look at the HW paper?!

- $n_i = 1,$
- $V_i = y_i^2,$
- $v = F_{x_i}^T \cdot y_i = \begin{pmatrix} f_1(x_i)y_i \\ \vdots \\ f_4(x_i)y_i \end{pmatrix}$
- $T_i = F_{x_i}^T * F_{x_i} = \begin{pmatrix} f_1(x_i)^2 & f_1(x_i)f_2(x_i) & \cdots & f_1(x_i)f_4(x_i) \\ \vdots & \vdots & \ddots & \vdots \\ f_4(x_i)f_1(x_i) & f_4(x_i)f_2(x_i) & \cdots & f_4(x_i)^2 \end{pmatrix}$

Oh... it is simple now and I know how to solve the task?! I need just implement four functions that calculates these variables n, v, V, T

```
In [1]: import numpy as np
import numpy.random as rnd

import matplotlib.pyplot as plt

%matplotlib inline
```

```
In [2]: WINDOW_START = 0
WINDOW_END = 4
WINDOW_STEP = 0.1

FIGURE_WIDTH = 20
FIGURE_HEIGHT= 10

M_SIZE = 4
A = np.array([1, 1, -1, 0.2])
```

So, let's define the canonical information. I thinks the tuple4 is the best solution, isn't it?

Canonical information in python is a tuple4 (n, v, V, T)

```
In [3]: def get_initial_cannonical_info(x, y):
    n = 1
    V = y * y
    Fxi = np.array([x ** k for k in range(M_SIZE)])
    v = Fxi * y
    T = np.matmul(np.expand_dims(Fxi, 1), np.expand_dims(Fxi, 1).T)

    return (n, V, v, T)
```

```
In [4]: def update_canonical_info(prev, x, y):
    n_prev, V_prev, v_prev, T_prev = prev
    n_curr, V_curr, v_curr, T_curr = get_initial_cannonical_info(x, y)

    n_new = n_prev + n_curr
    V_new = V_prev + V_curr
    v_new = v_prev + v_curr
    T_new = T_prev + T_curr

    return (n_new, V_new, v_new, T_new)
```

Let's try to draw a right function (red)


```
In [5]: def f(x):
    x_p = np.array([x ** k for k in range(M_SIZE)])
    return np.dot(A, x_p)
```

```
In [6]: def draw_right_function(plot):
    x = np.arange(WINDOW_START, WINDOW_END, WINDOW_STEP)

    y = [f(xi) for xi in x]

    plot.figure(figsize=(FIGURE_WIDTH, FIGURE_HEIGHT))
    plot.plot(x, y, 'r', label='f_real', linewidth=5)
```

```
In [7]: draw_right_function(plt)
```



Let's try to define estimate function

```
In [8]: def estimate_f(cannonical_info, x_i):
    #  $f(x) \hat{=} FxT^{-1}v$ 
    n_prev, V_prev, v_prev, T_prev = cannonical_info
    Fx = np.expand_dims(np.array([x_i ** k for k in range(M_SIZE)]), 0)
    T_1 = np.linalg.inv(T_prev)
    v = np.expand_dims(v_prev, 1)

    ## Please, don't ask me "What there happening, because I don't know.
    I found it from stackoverflow"
    return np.squeeze(np.matmul(np.matmul(Fx, T_1), v)).item()
```

Good. Next is the variance of the estimate function

```
In [9]: def variance_of_estimate_f(cannonical_info, x_i, eps):
    #  $Var(f(x)) \hat{=} \sigma^2 FxT^{-1}FTx$ 

    n_prev, V_prev, v_prev, T_prev = cannonical_info
    Fx = np.expand_dims(np.array([x_i ** k for k in range(M_SIZE)]), 0)
    T_1 = np.linalg.inv(T_prev)
    ## Please, don't ask me "What there happening, because I don't know.
    I found it from stackoverflow"
    return np.squeeze((eps * np.matmul(np.matmul(Fx, T_1), Fx.T))).item()
```

The last is monstrous function - Estimation of variance of the estimate functions... God damn!

```
In [10]: # $Var(f(x)) \hat{=} V - vTT^{-1}vn - m * FxT^{-1}FTx$ 
def my_monstrous_function(cannonical_info, x_i):
    n_prev, V_prev, v_prev, T_prev = cannonical_info

    Fx = np.expand_dims(np.array([x_i ** k for k in range(M_SIZE)]), 0)
    T_1 = np.linalg.inv(T_prev)

    denominator = n_prev - M_SIZE
    numerator = V_prev - np.squeeze(np.matmul(np.matmul(v_prev.T, T_1),
    v_prev)).item()
    multiplier = np.squeeze((np.matmul(np.matmul(Fx, T_1), Fx.T))).item()
    return (numerator / denominator) * multiplier
```

OK, Let's move on - make a sample with n = 10 and plot all the results

```
In [97]: def draw(plot, canonical_info, observations_x, eps):
    right_y = [f(x) for x in observations_x]
    estimated_y = np.array([estimate_f(canonical_info, x) for x in observations_x])

    var_y = np.array([variance_of_estimate_f(canonical_info, x, eps[idx])
]) for idx, x in enumerate(observations_x)])
    var_y_plus = estimated_y + var_y
    var_y_minus = estimated_y - var_y

    var_var_of_est_y = np.array([my_monstrous_function(canonical_info,
x) for x in observations_x])
    var_var_of_est_y_plus = estimated_y + var_var_of_est_y
    var_var_of_est_y_minus = estimated_y - var_var_of_est_y

    plot.figure(figsize=(FIGURE_WIDTH, FIGURE_HEIGHT))

    plot.plot(x, right_y, 'r', label = "Real", linewidth=5)
    plot.plot(x, estimated_y, 'b', label = "Estimate")
    plot.plot(x, var_y_plus, 'g', label = "Variance of Estimation")
    plot.plot(x, var_y_minus, 'g')

    plot.plot(x, var_var_of_est_y_plus, 'y', label = "Variance of variance
of Estimation")
    plot.plot(x, var_var_of_est_y_minus, 'y')

    plot.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespa
d=0.)
    plot.show()
```


```
In [102]: SAMPLE = 10
MEAN = 1
VAR = 0.5

MEAN_EPS = 0
VAR_EPS = 0.6

eps = rnd.normal(MEAN_EPS, VAR_EPS, SAMPLE)
x = rnd.normal(MEAN, VAR, SAMPLE)
x.sort() what for?
y = f(x) + eps

canonical_info = get_initial_canonical_info(x[0], y[0])
for xi, yi in zip(x, y):
    canonical_info = update_canonical_info(canonical_info, xi, yi)

draw(plt, canonical_info, x, eps)
```




```
In [105]: SAMPLE = 100
MEAN = 1
VAR = 0.5

MEAN_EPS = 0
VAR_EPS = 0.6

eps = rnd.normal(MEAN_EPS, VAR_EPS, SAMPLE)
x = rnd.normal(MEAN, VAR, SAMPLE)
x.sort()
y = f(x) + eps

canonical_info = get_initial_canonical_info(x[0], y[0])
for xi, yi in zip(x, y):
    canonical_info = update_canonical_info(canonical_info, xi, yi)

draw(plt, canonical_info, x, eps)
```



```
In [106]: SAMPLE = 1000
MEAN = 1
VAR = 0.5

MEAN_EPS = 0
VAR_EPS = 0.6

eps = rnd.normal(MEAN_EPS, VAR_EPS, SAMPLE)
x = rnd.normal(MEAN, VAR, SAMPLE)
x.sort()
y = f(x) + eps

canonical_info = get_initial_canonical_info(x[0], y[0])
for xi, yi in zip(x, y):
    canonical_info = update_canonical_info(canonical_info, xi, yi)

draw=plt, canonical_info, x, eps)
```

